

Speculative synchronization in StarSs, a task-based programming model

Rahulkumar Gayatri

Outline

- StarSs (sequential → parallel)
 - Compiler directives (pragmas)
 - Runtime features (Dependency Graph)
 - Synchronization directives
- Software Transactional Memory (STM) based concurrency control for critical memory updates
- Speculative task generation

- Moose Parallelization (current work)
 - OpenMP and Pthread parallelization
 - Optimizations

StarSs Programming Model

- Implicit parallel programming model for widely used multi-core architectures
- Single parallel application code for SMPs, Cell B.E, GPUs, cluster
- Sequential application with pragmas
 - Pragmas (compiler directives) to annotate a task
 - Runtime schedules the tasks on the available resources
 - Thread pool (Main and Worker threads)
 - Task Dependency Graph (TDG)

Task pragma and its clauses

Application code is divided into fine grained tasks

Task pragma and its clauses

Application code is divided into fine grained tasks

```
#pragma css task clauses  
void task (parameters)
```

Task pragma and its clauses

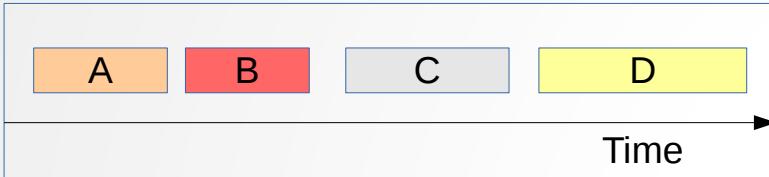
Application code is divided into fine grained tasks

```
#pragma css task clauses  
void task (parameters)
```

input (list of parameters)
inout (list of parameters)
output (list of parameters)

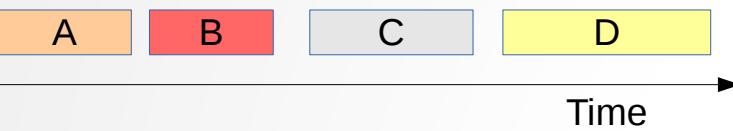
StarSs: annotation of tasks

Sequential code



StarSs: annotation of tasks

Sequential code



```
#pragma css task input(a) inout(b, c) \
output(d)

void A (a, b, c, d)

void B (e, f)

#pragma css task input(d) inout(b)
void C (d, b)

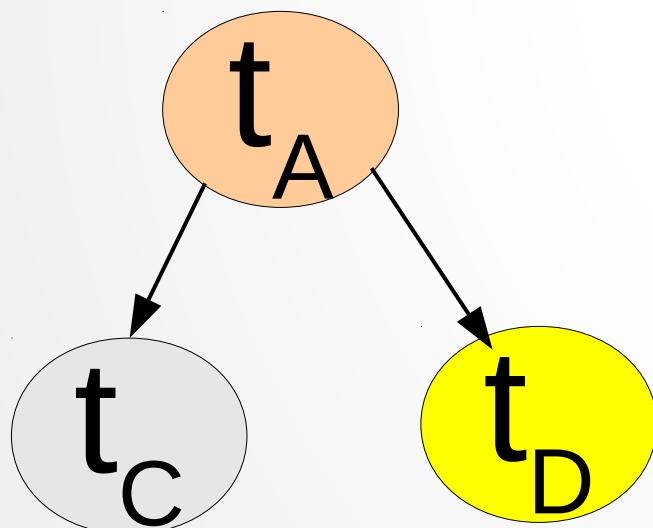
#pragma css task input(d) inout(c) output(e)
void D (d, c, e)
```

StarSs: TDG

Sequential code



TDG



```
#pragma css task input(a) inout(b, c) \
output(d)

void A (a, b, c, d)

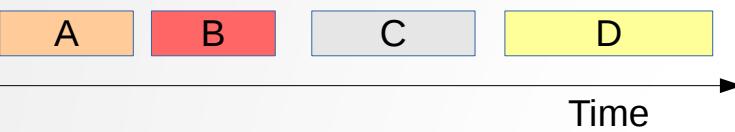
void B (e, f)

#pragma css task input(d) inout(b)
void C (d, b)

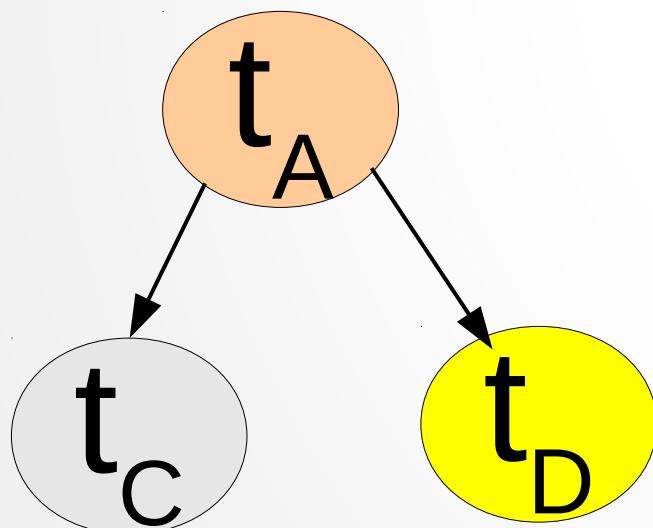
#pragma css task input(d) inout(c) output(e)
void D (d, c, e)
```

StarSs: scheduling

Sequential code



TDG



```
#pragma css task input(a) inout(b, c) \
output(d)

void A (a, b, c, d)

void B (e, f)

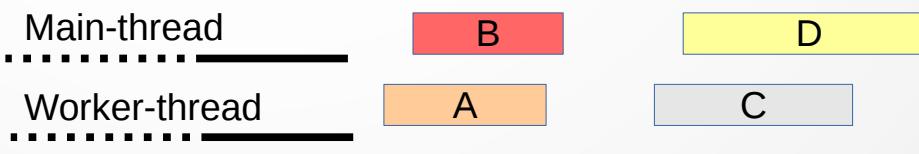
#pragma css task input(d) inout(b)

void C (d, b)

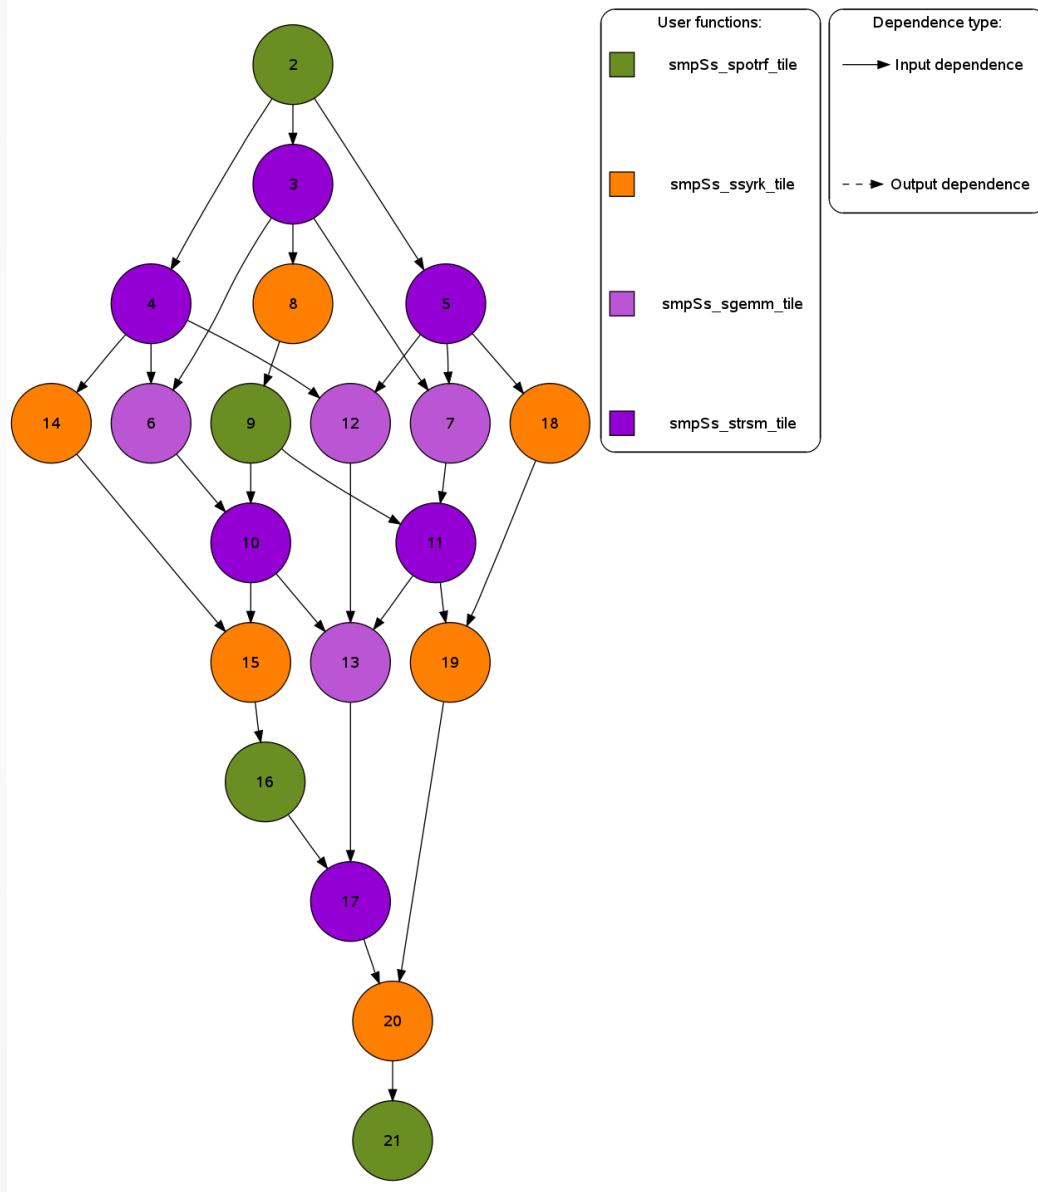
#pragma css task input(d) inout(c) output(e)

void D (d, c, e)
```

SMPSSs scheduling



Cholesky TDG



StarSs: critical updates

- Reduction clause
 - does not create a dependency edge
- Need synchronization to maintain correctness

StarSs: locks

```
# pragma css task input (n, j, a[n]) inout (results) \
                    reduction (results)

void nqueens_task (int n, int j, char* a, int* results)
{
    ...
# pragma css mutex lock (results)
    *results = *results + local_sols;
# pragma css mutex unlock (results)
    ...
}
```

Drawbacks of lock-based synchronization

- Deadlock – two concurrent threads wait for each other to release a resource, thus neither progress
- Livelock – state of the threads change with regard to each other but neither progress
- Priority inversion – Lower priority thread blocks the higher priority thread
- Difficult to compose and debug
- Not scalable

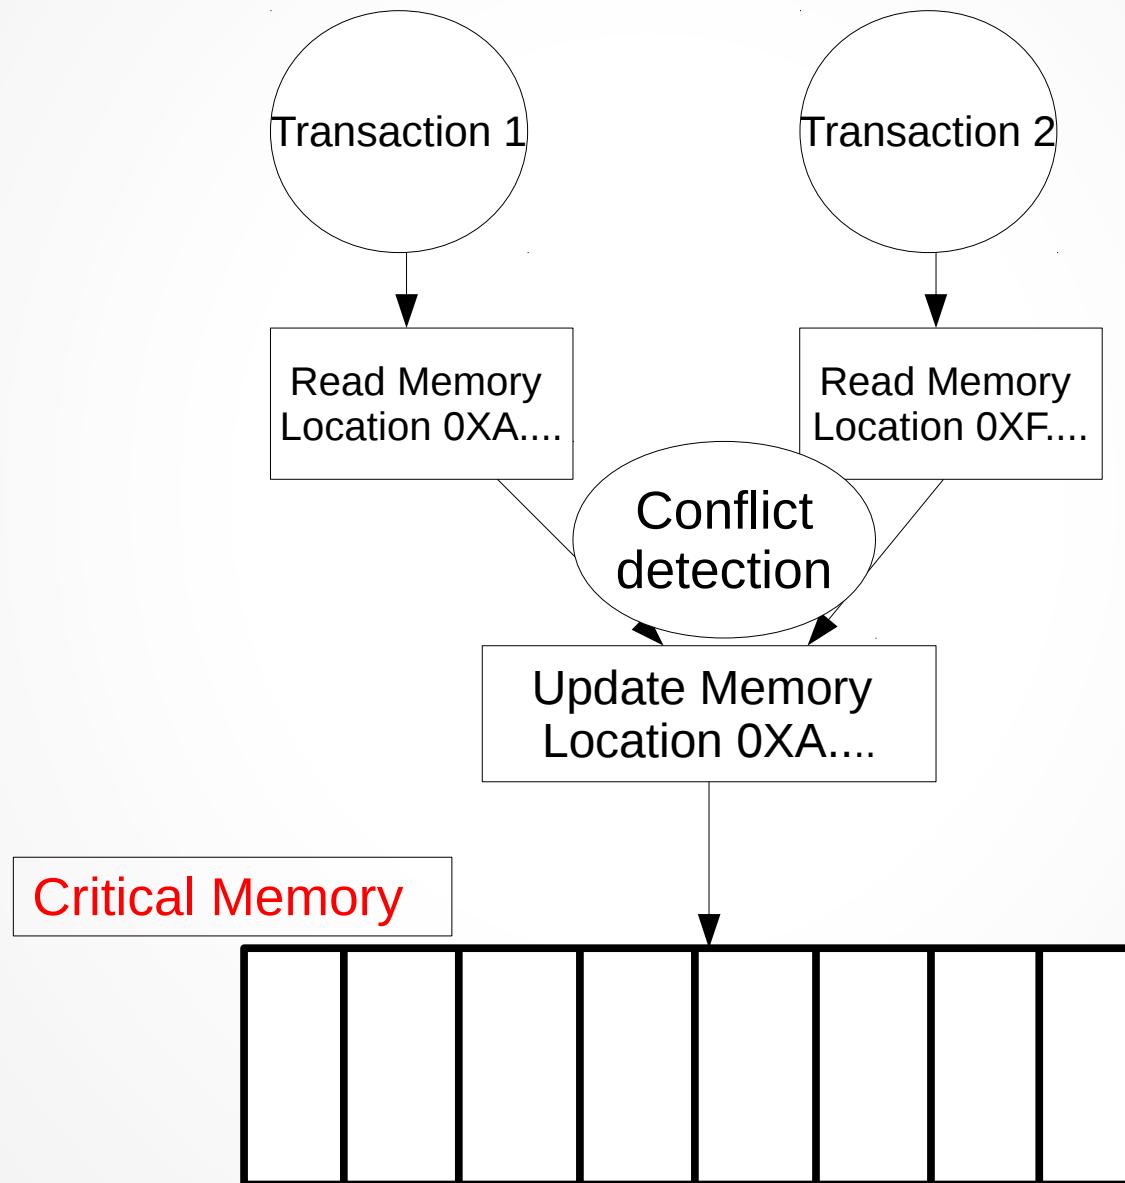
Software Transactional Memory (STM)

Transaction

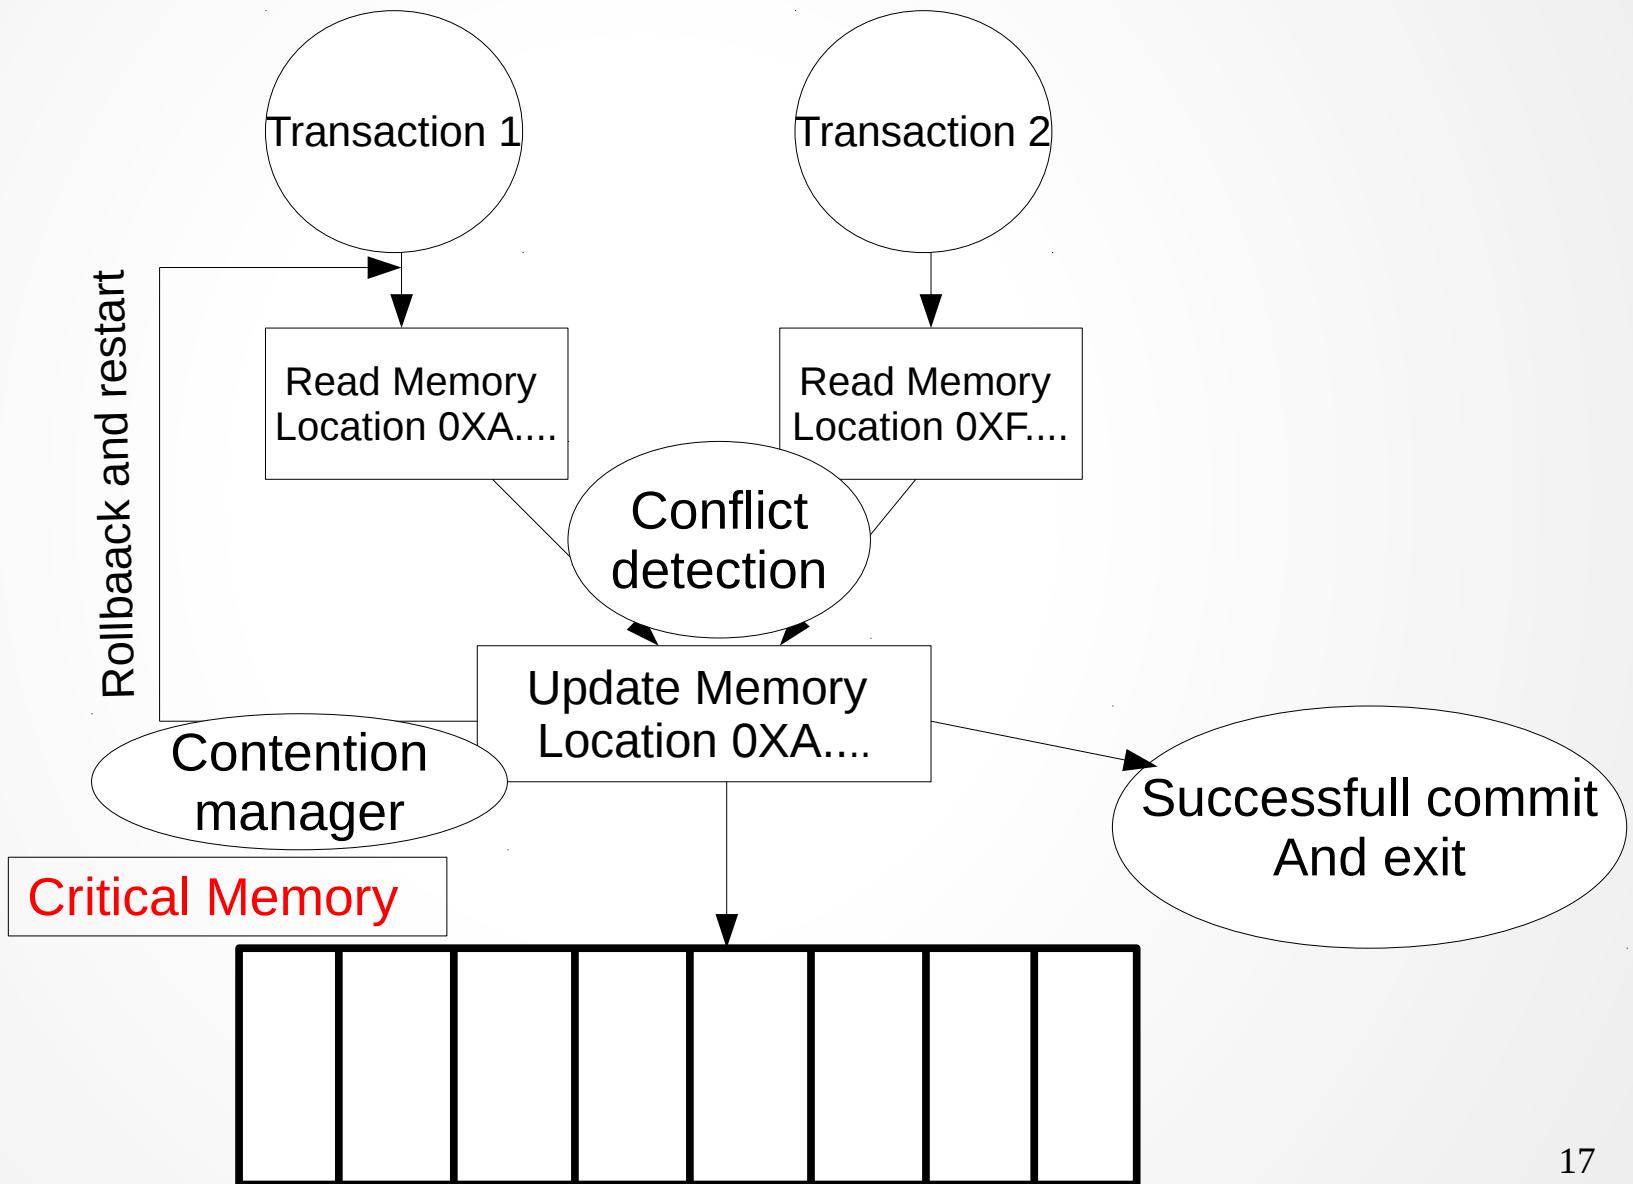
Series of read
and writes to
memory



STM: conflict detection



STM: contention manager



TinySTM

- A light weight STM library
- 3 versions
 - WRITE_BACK_ETL (check conflicts at encounter time)
 - WRITE_BACK_CTL (check conflicts at commit time)
 - WRITE_THROUGH (directly update shared memory)

TinySTM in StarSs

- Main and worker threads perform the necessary implementations
 - Changes in configuration for the appropriate linking
- Lock pragma starts a transaction and unlock pragma commits the transaction

Compiler vs runtime

- Stack calls in TinySTM library make runtime implementation difficult
- StarSs compiler replaces lock and unlock pragmas with TinySTM transaction calls.

Compiler transformations of lock and unlock

Application code

```
...
#pragma css mutex lock (results)
    *results = *results + local_sols ;
#pragma css mutex unlock (results)
...
*results = *results + local_sols ;
```

Start the transaction

Save the stack context

Load parameter into the transactional context

Store the updated parameter to memory
Commit the transaction

Applications with reductions

- Nqueens
 - BFS
 - Specfem3D
 - Gmeans
-
- The diagram consists of two vertical lines. The left line connects 'Nqueens', 'BFS', and 'Specfem3D' to the right text 'STM better than Locks'. The right line connects 'Gmeans' to the right text 'Locks better than STM'. Each line has a vertical segment on its left and a horizontal arrow pointing to the right text.
- STM better than Locks
- Locks better than STM

Locks vs Transactions performance

- Both versions scale similarly
- STM performed better in applications with high lock contention
- More beneficial to update multiple memory locations in a single transaction

STM overhead

- STM is notorious for high software overhead
 - Setting up transaction environment
 - Saving stack information
 - Conflict detection
 - Contention manager
- Hybrid-TM is the way forward

2nd Phase of PhD

Speculative memory updates to speculative task execution

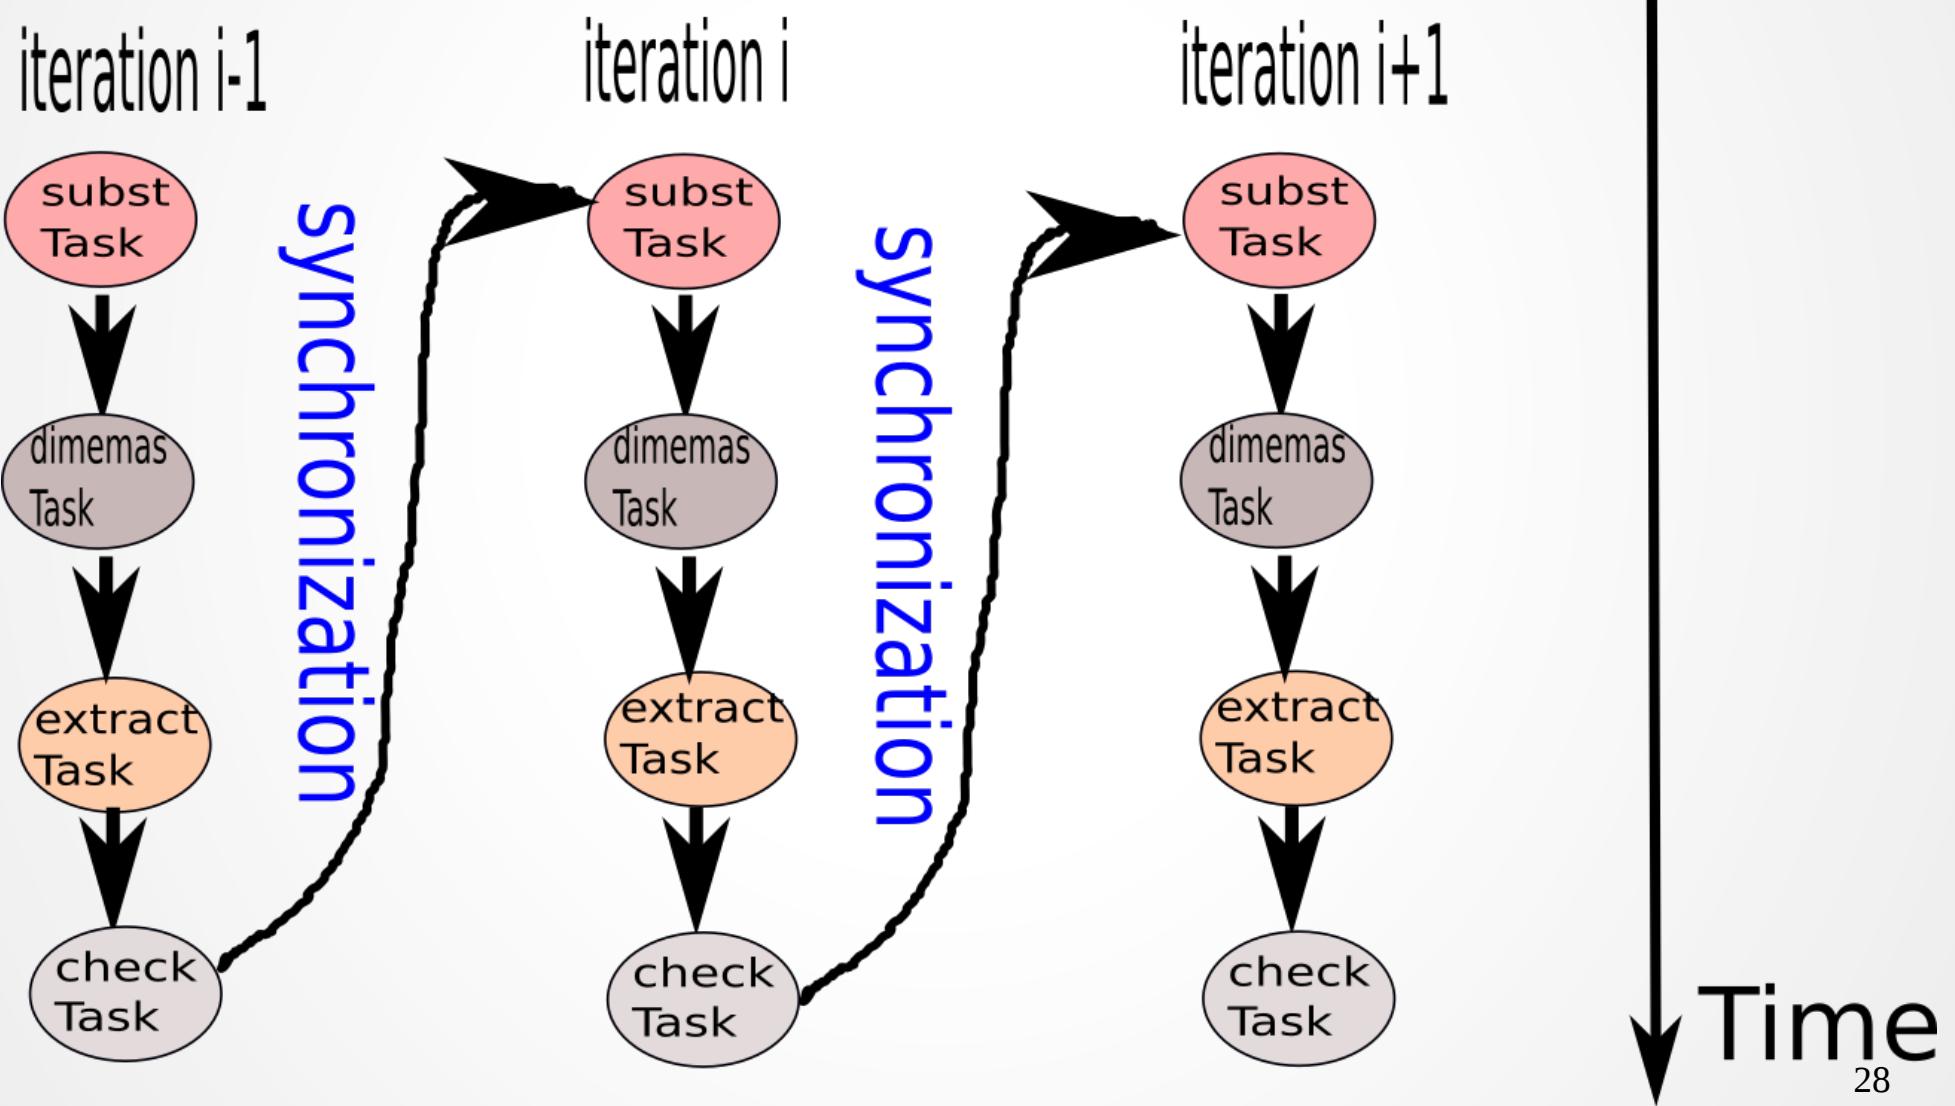
Necessary of task synchronization

- Separation of task generation and execution makes synchronization a necessity
 - #pragma css **wait on** (addr)
 - #pragma css **barrier**

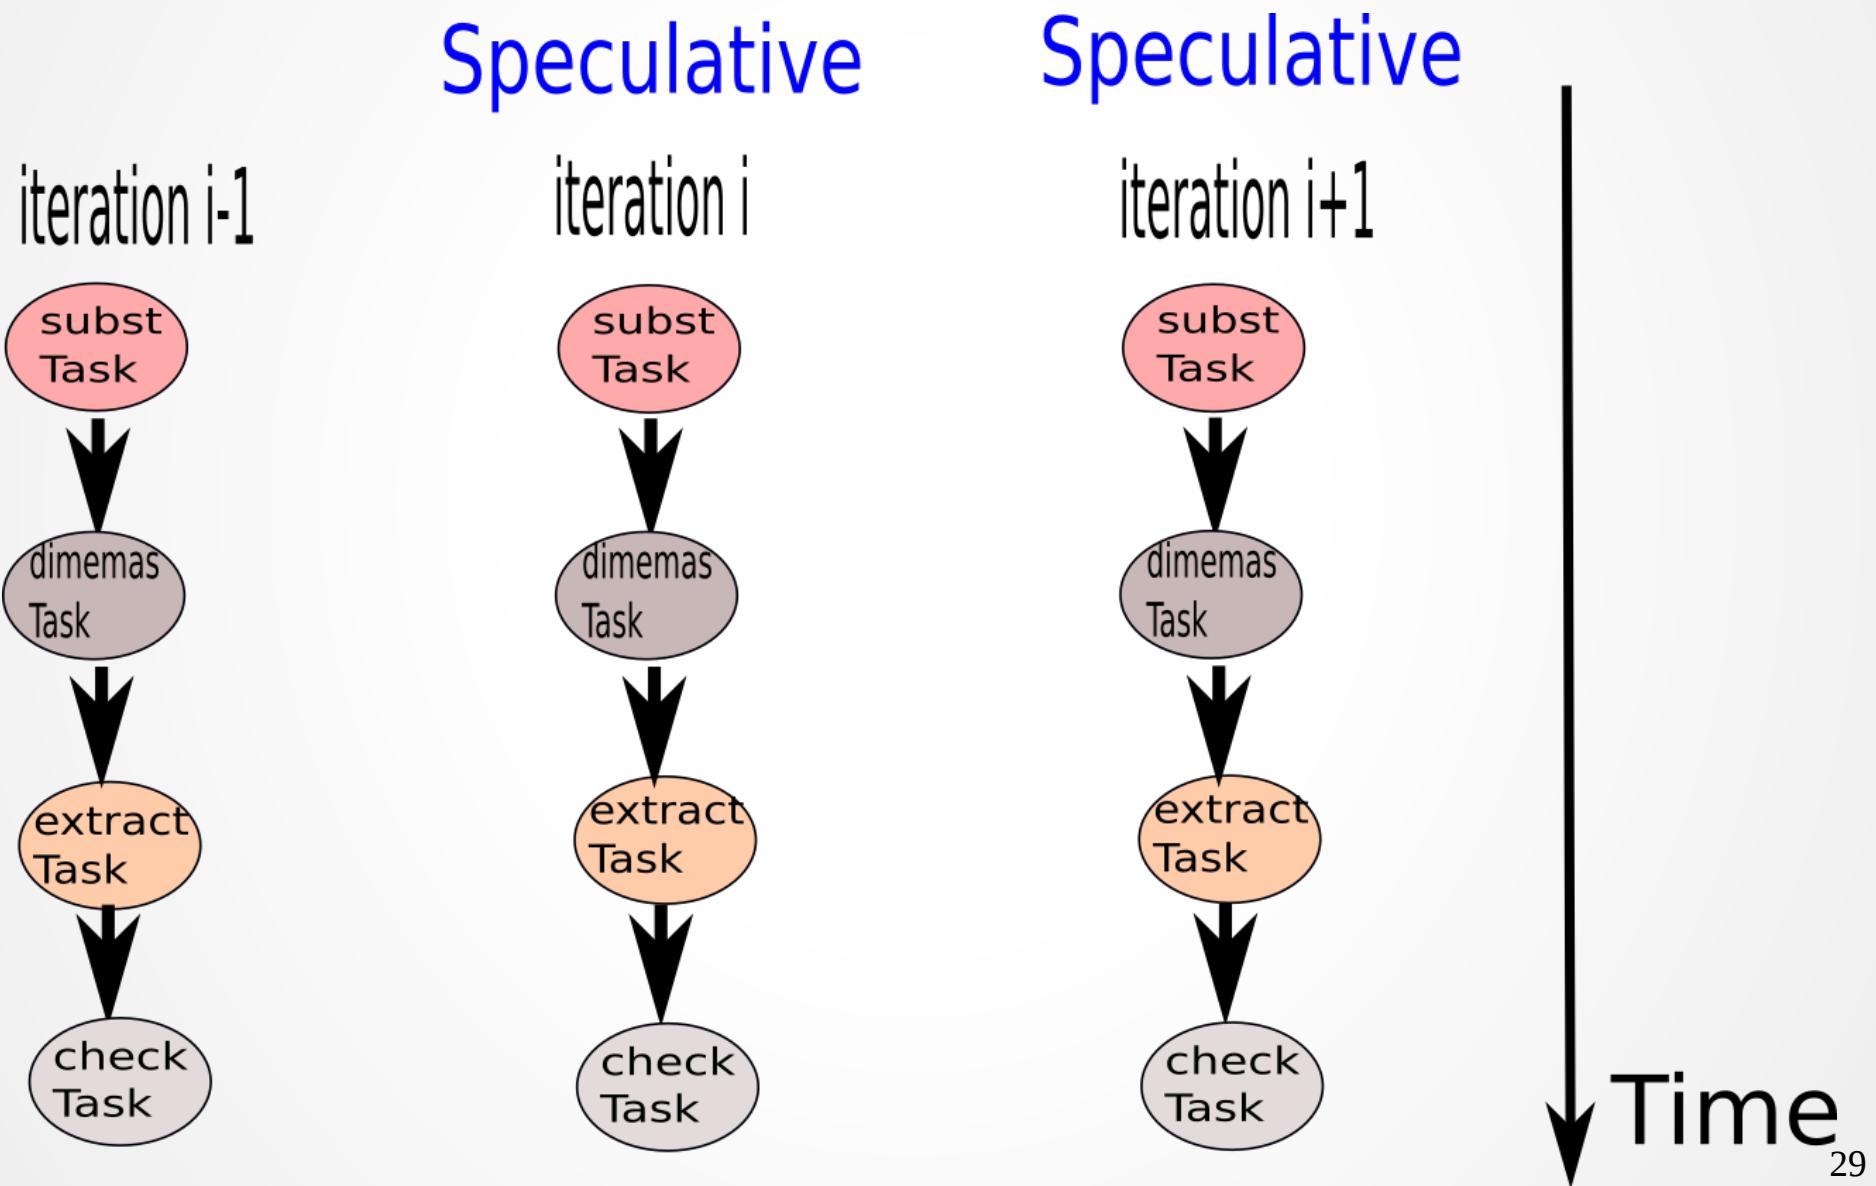
Usage of synchronization

```
while(!goal_achieved)
{
    newBFD = pop_queue();
    subst(refCFG, newBFD, newCFG);
    dimemas(newCFG, trace, dimOUT);
    extract(newBWD, dimOUT, finalOUT);
    check(finalOUT, goal_achieved);
    #pragma css wait on(goal_achieved)
}
```

TDG



Speculative TDG (ideal)



Speculative directive: clauses

```
#pragma css speculate values(finalOUT) wait(goal_achieved)

while(!goal_achieved)

{
    newBFD = pop_queue();
    subst(refCFG, newBFD, newCFG);
    dimemas(newCFG, trace, dimOUT);
    extract(newBWD, dimOUT, finalOUT);
    check(finalOUT, goal_achieved);
    #pragma css wait on(goal_achieved)
}
```

Speculative directive: clauses

```
#pragma css speculate values(finalOUT) wait(goal_achieved)

while(!goal_achieved)

{
    newBFD = pop_queue();
    subst(refCFG, newBFD, newCFG);
    dimemas(newCFG, trace, dimOUT);
    extract(newBWD, dimOUT, finalOUT);
    check(finalOUT, goal_achieved);
}
```

values – protection in case of mis-speculation
wait – determines the progress of the loop

Success of Speculation

- Valid tasks – Successfull speculation
 - Tasks generated even with the synchronization
- Invalid tasks – Undo their results
 - Loop predicate evaluated to false

Speculative directive: implementation

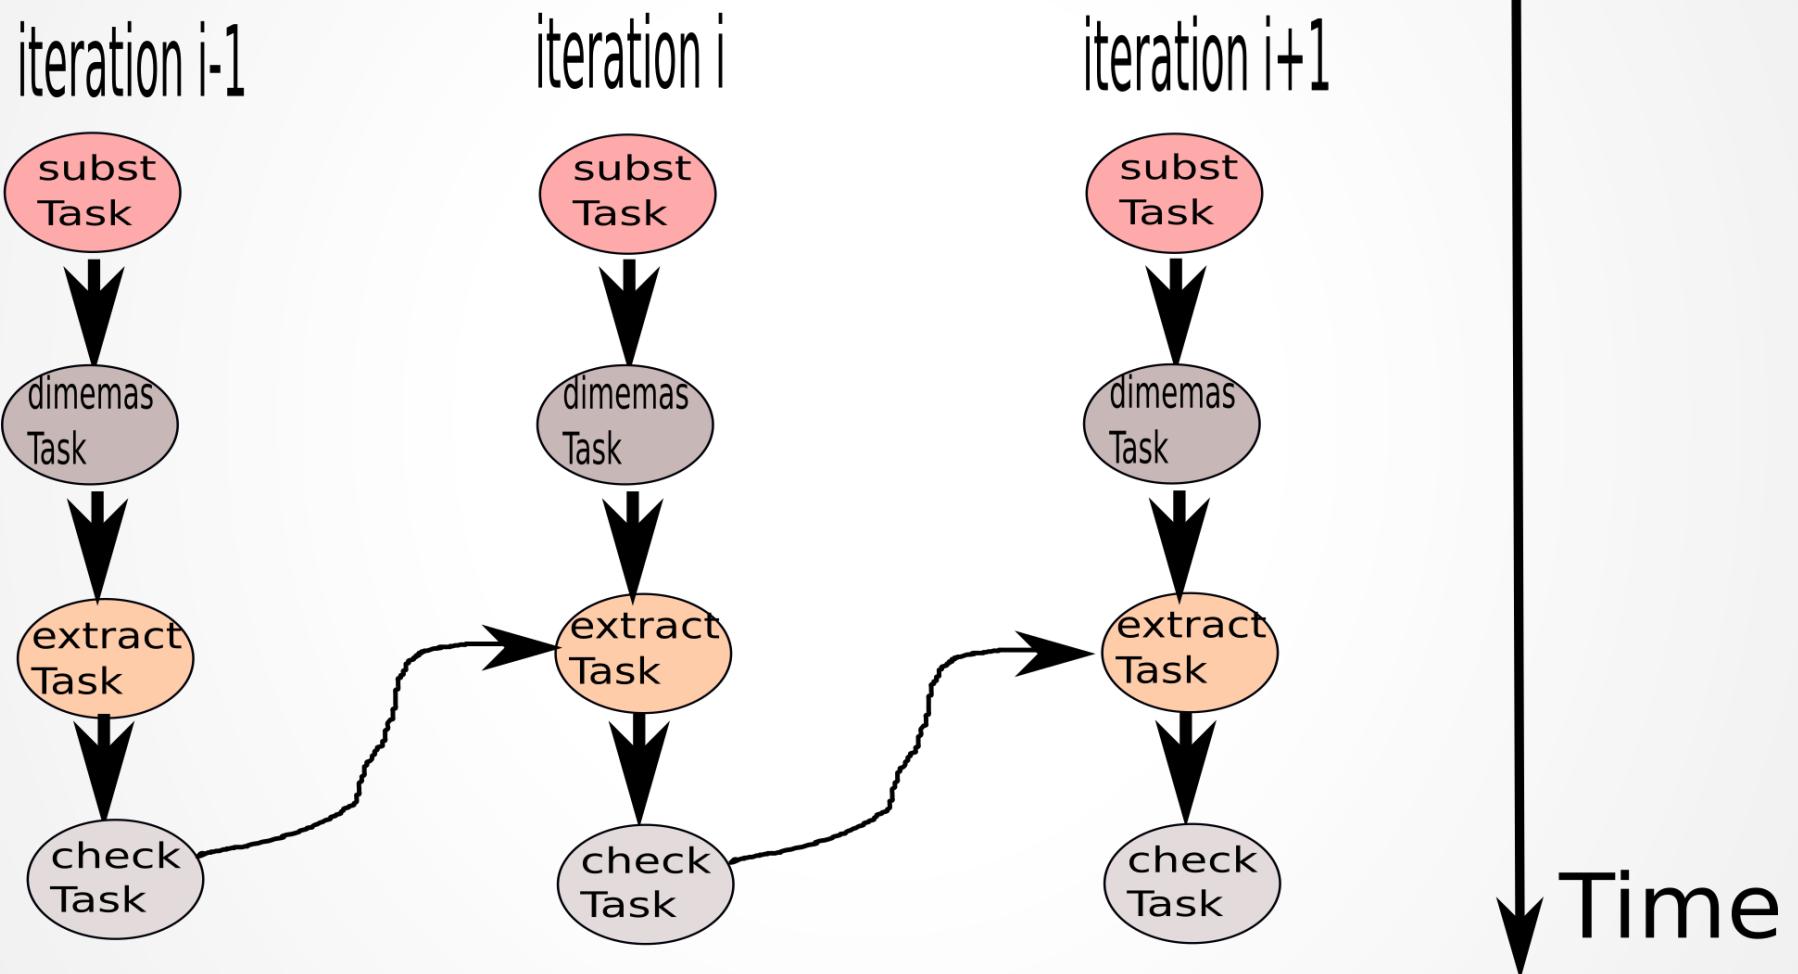
- *Speculative* tasks - transactions
- Parameters from **values** are loaded into the transactional context
- Store values only if speculation is successful
- Use of a **guard** function to evaluate the success of speculation

Guard function

Guard function

```
int guard (void *speculate_params[1])  
{  
    return (!*(speculate_params[0]));  
}  
speculate_params[0] = &goal_achieved ;
```

Speculative TDG



Types of Speculative tasks

- Data-dependent speculative tasks
 - Cross-iteration dependencies
 - Tasks from iteration i consume values from tasks in $i-1$
 - Overlap task generation with task execution
- Control dependent speculative tasks
 - Tasks from iteration i are independent of tasks from $i-1$
 - Parallel execution of tasks

Applications for testing

- Gauss – Seidel
- Jacobi
- Red-black
- Kmeans

Observations

- Non-speculative versions perform better
 - Both versions scale similarly
- Better performance with increasing number of threads
- TinySTM overhead cannot be overcome with the additional parallelism obtained

TinySTM overhead

- Less number of bigger transactions is beneficial
 - Bigger tasks more beneficial than smaller tasks
 - Longer rollbacks versus higher number of TinySTM calls
- Increase task granularity
 - Instead of increasing the number of tasks in the TDG, we reduce them to compensate for the overhead

Data-version based speculative task execution

- TinySTM load and store a major overhead
- Redundant conflict detection mechanisms
- Data-versions to maintain correctness

Data-version based speculative task execution

- TinySTM load and store a major overhead
- Redundant conflict detection mechanisms
- Data-versions to maintain correctness

Overhead

- 1) 1 copy for successful speculation
- 2) 2 copies for mis-speculation
- 3) Evaluation of the **guard** function

Observations

- On average a 20% performance improvement over non-speculative application
 - Significant since the improvement is over an already parallelized code
- The benefits of speculation are better observed with increasing number of threads
- With increasing problem sizes more number of threads are needed for performance benefits

Optimizations

- Update original memory locations inside speculative tasks
- Limit on the number of iterations to be speculated upon
 - `speculation_tasks`, an environment variable

Current Work

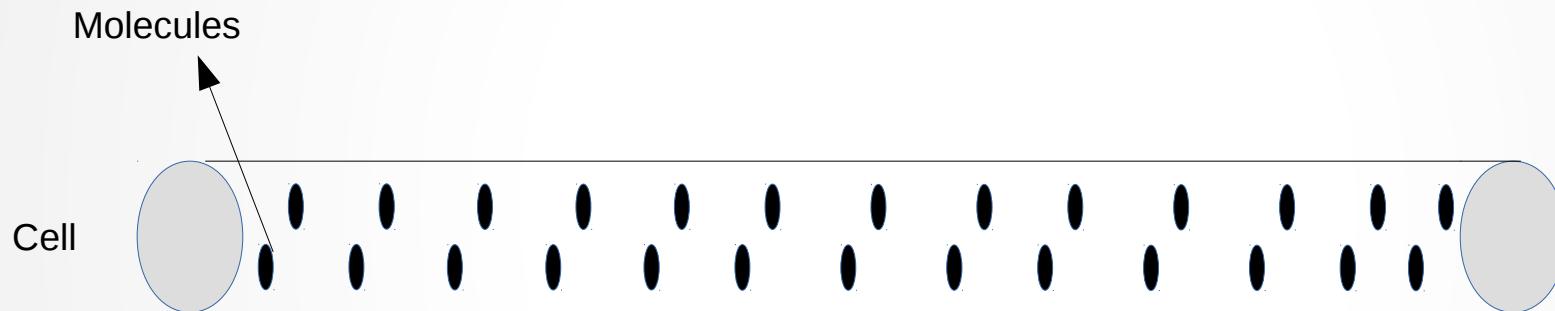
Parallelization of Moose

Current Work

MOOSE, a Multiscale Object-Oriented
Simulation Environment.

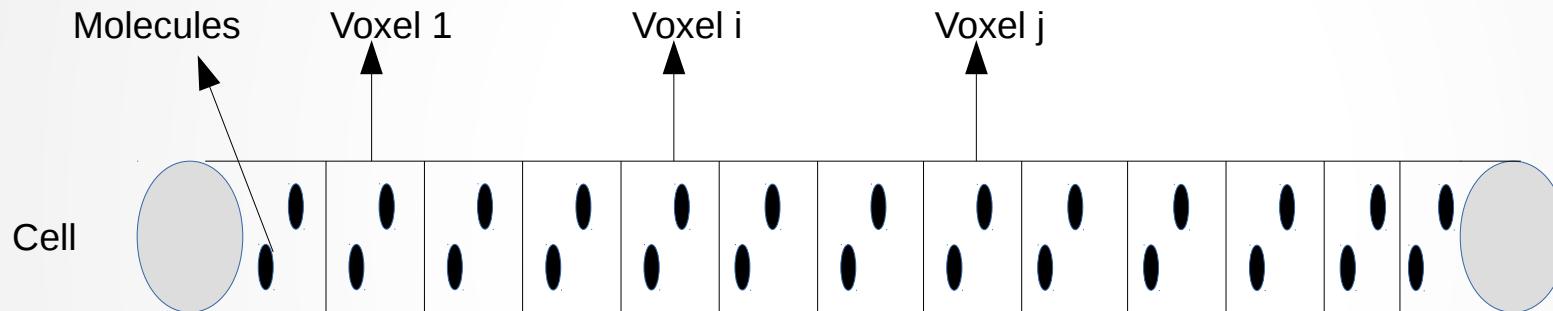
Cell/Neuron division

MOOSE, a Multiscale Object-Oriented
Simulation Environment.



Cell/Neuron division

MOOSE, a Multiscale Object-Oriented
Simulation Environment.



- A cell is subdivided into Voxels
- Independent computations in each voxel at a given timestep

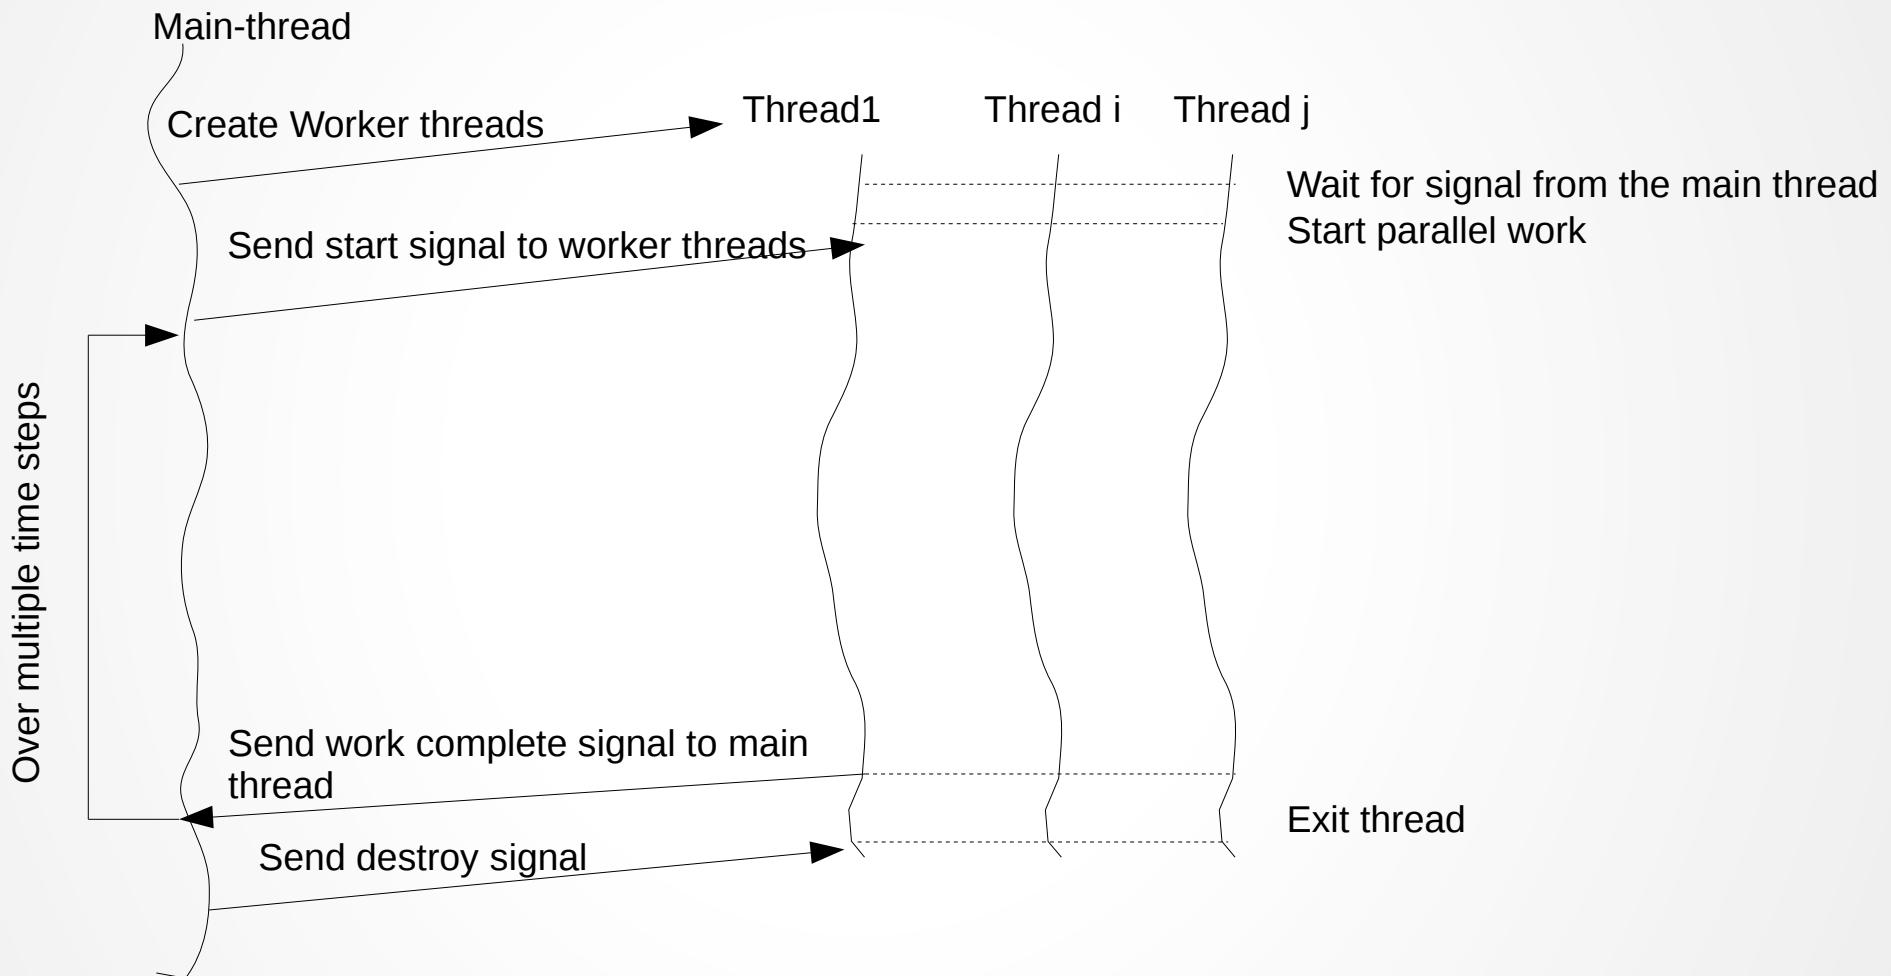
Single cell simulation

- Chemical interactions on the molecules
- Kinetic solver
 - Large volume (deterministic behavior)
 - Simulates the behavior of the molecule using linear transformation (rungekutta method of order 5)
- Stochastic solver
 - Small volume
 - Simulation is performed over individual molecules
- Diffusion solver
 - Calculate the diffusion gradient

Programming Models used

- OpenMP
 - Parallel-for and task-based parallelization
 - Ease of use
- Pthreads
 - More flexibility
 - Optimized on thread creation-destruction by the use of semaphores

Optimization of pthreads using semaphores



Optimizations

- Memory alignment
- Reduce the number of OpenMP parallel pragmas
- Remove hyperthreading

Results of Parallelization on 8 threads

- Kinetic solver
 - OpenMP – 5X
 - Pthreads – 4.5X
- Stochastic Solver
 - OpenMP – 5.5X
 - Pthreads – 4.8X
- Diffusion solver
 - OpenMP – 3.8X
 - Pthreads – 3.7X

Conclusion

- PhD Thesis
 - Compiler and runtime development
 - Speculative synchronization (using STM)
- Cell simulation
 - Parallelization of solvers for chemical interactions

THANK YOU!!